

• Process Concept

An OS runs many programs :-

1. Batch System : Handles tasks
2. Time-shared System : Handles user / prog tasks

A process has diff parts :-

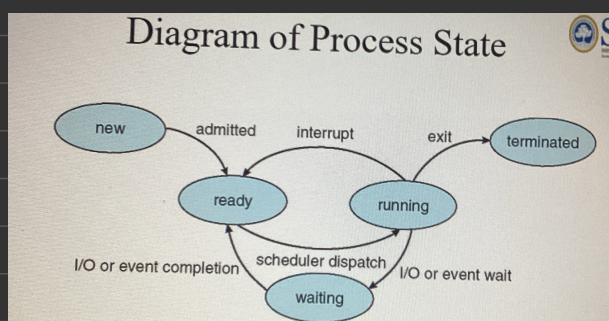
- i) code : The actual prog instruc
- ii) PC & reg : track where the process is there in execution
- iii) Stack : Stores temporary data
- iv) data section : stores global vars
- v) Heap : store dynamic alloc memory

Also if prog is **passive**, it's stored on disk and if **active** then run in memory

Process State :-

As prog. runs it moves through diff states

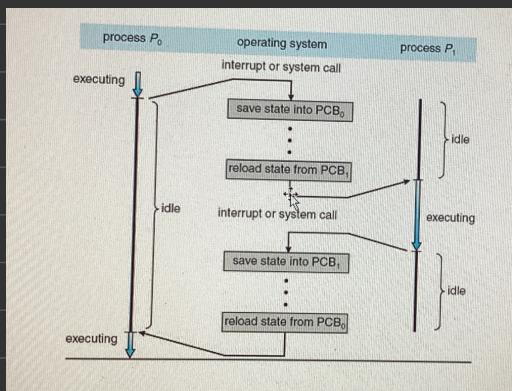
1. New : It's created
2. Running : Currently executing
3. Waiting : waits for some event / input
4. Ready : ready to run but waits for its turn in CPU
5. Terminate : done



- Process Control Block (PCB)

every process has a pcb that stores :-

1. State : running, ready , etc
2. PC : location of instruc to execute next
3. CPU reg : have info specific to the process
4. CPU scheduling : checks priorities & queue pointers
5. Memory mgmt info : how much memory process uses
6. Accounting info : CPU used , clock time elapsed since start
7. I/O status info : which I/O process used

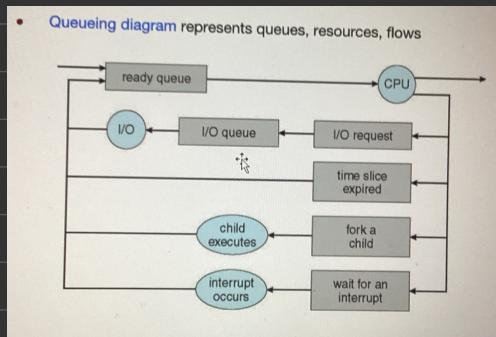


• Process Scheduling

OS uses scheduling to maximize CPU use by quickly switching between processes

There are diff queues :-

1. Job queue : set of all processes in the system
2. Ready queue : processes ready to run
3. Device queues : processes waiting for I/O devices



Schedulers :-

1. Short - term scheduler (CPU scheduler) :-

- Selects which process will run next
- It's invoked frequently

2. Long-term Scheduler (Job Scheduler) :-

- decides which process gets added to system
- Processes are I/O bound & CPU bound
- Good mix of both

3. Mid-term Scheduler

- Can pause processes if too many
- Swap them
- Bring again when needed

Note * Context Switch: CPU switches processes, saves the old one & loads new one, no real work done here

- Process Ops
- Process Creation

1. A parent process creates child process
2. Processes can share resources with child or their own
3. Par & child share resources (may/may not)
4. Both execute concurrently

How to create : `fork()` : new process

`exec()`: loads a new prog in process

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
pid_t pid;

/* fork a child process */
pid = fork();

if (pid < 0) { /* error occurred */
    fprintf(stderr, "Fork Failed");
    return 1;
}
else if (pid == 0) { /* child process */
    execvp("/bin/ls", "ls", NULL);
}
else { /* parent process */
    /* parent will wait for the child to complete
    wait(NULL);
    printf("Child Complete");
}

return 0;
}
```

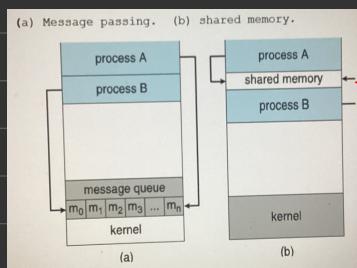
• Process termination

1. process calls `exit()` system call
2. parent can also terminate child process by `abort()`
3. If parent waits for child it'll use `wait()`
4. If not then child becomes **zombie**
5. If parent terminates & child runs, he's an **orphan**

• Communication (Interprocess Comm IPC)

- Processes can be independent or cooperating
- Cooperating process have benefits like info share, fast computation, convenience & modularity
- Cooperating needs IPC
- Two models of IPC :-

1. IPC - Shared memory



- processes share a specific area for comm.
- The comm is under control of user & not OS
- It lacks sync to avoid data inconsistencies when multiple processes access shared memory

2. IPC - Message passing

- processes comm. by sending & receiving msg
- has 2 ops :-
 - send (msg)
 - receive (msg)
- Also processes need to establish a comm link to exchange msgs
- Implementation of comm link
 - ① physical : Shared memory
h/w bus
Network
 - ② logical : direct / indirect
Sync / Async
Automatic

- Direct / Indirect Comm

1. Direct :-

- processes comm. by naming each other
eg send (P, msg)
- Some properties of comm link :-
 - i) established automatically
 - ii) bidirectional
 - iii) only b/w two process

2. Indirect :-

- Msgs sent & received through mailbox
- each mailbox has unique id
- Here processes only comm via mailbox
- props :-
 - i) uni / bi-dir
 - ii) each pair may have multiple comm links
- ops are :-
 $\text{send } (A, \text{msg})$ A: Mailbx
 $\text{receive } (A, \text{msg})$

Note * Blocking (Sync) : waits till msgs sent / received
non-blocking (osync) : doesn't

Buffering in Message Passing

Messages are buffered and queued:

- Zero Capacity: No queue, processes wait for immediate message handling.
- Bounded Capacity: Limited buffer space, sender waits if the buffer is full.
- Unbounded Capacity: Infinite queue, sender never waits.

• Eg's of IPC system

1. POSIX

- For shared memory, segments are created using `shm_open()`
- memory size set using `ftruncate()`
- To print `sprintf()`

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* strings written to shared memory */
    const char *message_0 = "Hello!";
    const char *message_1 = "World!";

    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* create the shared memory object */
    shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);

    /* configure the size of the shared memory object */
    ftruncate(shm_fd, SIZE);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_WRITE, MAP_SHARED, shm_fd, 0);

    /* write to the shared memory object */
    sprintf(ptr, "%s", message_0);
    ptr += strlen(message_0);
    sprintf(ptr, "%s", message_1);
    ptr += strlen(message_1);

    return 0;
}
```

```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* open the shared memory object */
    shm_fd = shm_open(name, O_RDONLY, 0666);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_READ, MAP_SHARED, shm_fd, 0);

    /* read from the shared memory object */
    printf("%s", (char *)ptr);

    /* remove the shared memory object */
    shm_unlink(name);

    return 0;
}
```

2. Mach

- mach uses msg-based comm
- Sys calls ~ msg
- each task get mailbox
- Sys calls are `msg_send`, `msg_receive`, `msg_rpc`
- It has flexible handling like wait, return, etc

3. Windows :-

- msg passing occurs via Local procedure calls
- procedure does ports like mailboxes
- Server creates private ports
- Client opens connection

• Client Server Systems

1. Sockets
2. Remote Procedure Calls
3. Pipes
4. Remote Method Invocation

• Sockets

- They are network comm. endpoints
- It's a endpoint for com
- defined by an IP add & a port
- Allows two-way comm
- * all ports <1024 are well-known
- Sockets in java :-
 - i) TCP (connection-oriented)
 - ii) UDP (connectionless)
 - iii) Multicast - send data to multiple users

• RPC

- They abstract function calls b/w processes on diff systems
- Stubs : proxy that does comm b/w client & server
- windows uses MIDL (Microsoft Interface definition lang)
- data representation in Big-endian & little-endian

• Pipes

- Allows processes to comm by passing data b/w them

- Ordinary pipes :-
 - i) They can't be accessed outside process
 - ii) usually b/w parent child
 - iii) unidir
 - iv) producer at one end & consumer at other

- Named pipes :-
 - i) More powerful than ordinary
 - ii) No parent-child relationship
 - iii) provided on both UNIX & windows
 - iv) Bi-dir

• Threads

- Threads are basic units of CPU
- They allow multiple task to run parallel
- easier & light
- simplify coding
- Improve efficiency
- Kernel OS is multithreaded



• Multicore programming

They put pressure on programmers & challenges like :-

1. divide work
2. balancing workload
3. split data
4. merge data
5. Test & debug

parallelism : multiple task at same time

Concurrency : making progress on multiple task not at same time

data & task parallelism

\downarrow \downarrow

distributed subsets of same data across multiple cores	distributes threads across cores
---	-------------------------------------

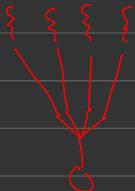
Note * Amdahl's law : predicts performance by adding CPU cores as more cores add up , speed limit

$$\text{Speedup} \leq \frac{1}{S + (1-S)/N}$$

• Multithreading models

1. Many - to - one

- Multiple user threads mapped to one kernel thread
- If one blocked, all blocked
- No parallelism
- eg: Solaris green threads



2. One - to - One Model

- Each user thread mapped to a kernel thread
- More concurrency
- Has overhead limits
- eg Windows, linux, Solaris 9+

3. Many - to - Many

- Allows many user to many many K threads
- More flexible & parallelism
- Keeps K threads sufficient
- eg Windows

4. Two - level

- like many - to - many , but allows binding a specific user to K threads
- eg IRIX , HP-UX , Solaris 8 , Tru64 UNIX
- Threading issues :-

1. Semantics of fork() & exec()

- does fork() duplicate calling thread or all ?
- exec() replaces curr thread

2. Signal Handling

- They manage signals
- every signal has a default handler or user-defined
- For multi-thread signals are delivered to specific or all threads

3. Thread Cancellation

- They can be cancelled immediately or after check
- Thread to be "is target thra
- In deferred, they check specific points where cancellation needed

async

deferred

4. Thread local storage

- Allows each thread to have its own copy of data
- diff from local vars or static data

5. Schedule activation &

- Mech to manage thread in models
- use an intermediate structure called lightweight process (Lwp)

• Synchronisation

- ↳ Topics :
 - 1. The critical section problem
 - (imp only) 2. Peterson's sol"
 - 3. Mutex locks

• Basics

- Sync make sure multiple proc work w/o interfering
- To maintain data consistency
- To avoid interruption & data issues

eg producer code : if space is there then add & inc (++)

```
while (true) {
    /* produce an item in next produced */
    while (counter == BUFFER_SIZE) {
        /* do nothing */
        buffer[in] = next_produced;
        in = (in + 1) % BUFFER_SIZE;
        counter++;
    }
}
```

consumer : if there are items , takes one > dec (--)

```
while (true) {
    while (counter == 0)
        /* do nothing */
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    counter--;
    /* consume the item in next consumed */
}
```

1. Critical section problem

- It's a section of code where a process accesses shared data
- Only 1 process can be in its crit sectⁿ at a time
- The objective to create protocols to access only 1 process

Algorithm : do {
 while (turn == j);
 // critical //
 turn = j ;
 // remainder //
} while (true)

Basically wait
till turn comes

Solution :-

- Mutual exclusion: only 1 in crit sectⁿ
- Progress : if no process in " " then take next process without any delay
- Bounded waiting : limit to process of no. of times in crit sectⁿ
- Kernel handles thin 2 ways :-
 1. pre-emptive : allows pre-emption
 2. Non-pre : exits kernel , then allows

2. Peterson's Solution

- It's a two-process solⁿ
- It's a good algo to solve crit section for 2 processes
- The 2 processes share 2 vars and turn up the flag of whose turn it is

Algorithm : do {
 flag [i] = true ;
 turn = j ;
 while (flag [j] & & turn == j);
 // critical //
 flag [i] = false ;
 // remainder //
} while (true)

Solution :-

- Mutual exclusion preserved : A proc only enters where the other one is not ready
- progress guaranteed : if no one in crit sectⁿ, anyone can enter
- Bounded limit

3. Mutex locks

- Simplifies the crit section problem
- It requires a lock before entering & release it afterward
- Functions are `acquire()` & `release()`

Algorithm :

```
acquire () {  
    while (! available)  
        // busy , plz wait //  
    available = false ;  
}  
release () {  
    available = true ;  
}  
do {  
    acquire lock ;  
    // crit sec //  
    release lock ;  
    // remainder //  
} while (true) ;
```

• Semaphores

- It's a tool for process sync to help em coordinate their actions better than mutex locks

- eg `Semaphore S` ;

two ops `wait()` & `signal()`

/ \

dec sema inc sema
val (if 0 val
then waits)

- Sema usage / implementation

1. Counting Sema : can have wide range values
2. Binary Sema : 0 / 1 val

- Ensure no 2 processes in wait() & signal() i.e. these both need to be in crit sect

- Implementation w/o busy waiting

- each sema has waiting queue
- each entry tracks sema's val & pointer next to it
- ops are just add a process in queue & removes & allow it to run

```
algo : wait (semaphore *s) {  
    s->val --;  
    if (s->val < 0) {  
        add to s->list;  
        block();  
    }  
}
```

```
signal ( semaphore *s ) {  
    s->val ++;  
    if (s->val <= 0) {  
        remove from s->list;  
        wakeup(p);  
    }  
}
```

- Deadlock & Starvation

- Deadlock : When 2 prcs wait forever for events that can be done by one of them

- Starvation : A prc may never run & just wait

- priority inversion : low-prior has lock of high-prior

- problems in sync

1. Bounder-buffer
2. Readers - writers
3. Dining philo

• Bounded - Buffer

- A fix no. of buffers can hold items
- In semq :-
 mutex : To control access to buffer
 full : Counts items in buffer
 empty : counts empty slots

Algo : For producer :-

```
do {
  wait (empty);
  wait (mutex);
  signal (mutex);
  signal (full);
} while (true);
```

For consumer :-

```
do {
  wait (full);
  wait (mutex);
  signal (mutex);
  signal (empty);
} while (true);
```

• Readers - writer Problem

- Multiple process share data
- Reader only read , not write
- Writers do both
- Goal is to allow multiple readers & one writer

Algorithm : Writers :-

```
do {
  wait (rw_mutex);
  signal ( " " );
} while (true);
```

Readers :-

```
do {
  wait (mutex);
  read_count++;
  if ( " " == 1 )
    wait (rw_mutex);
    signal ( mutex );
  wait (mutex);
  read_count--;
  if ( " " == 0 )
    signal ( rw_mutex );
    signal ( mutex );
} while (true);
```

Note : 1. no reader shld wait when writer uses process
 2. writer shld write asap

(it also leads to) ↗ both starvation

• Dining - Philosophers problem

- It leads to deadlock
- prevention :-
 1. limit no. of philosophers
 2. Allow philo to pick a chopstick in a specific order
 3. Allow atmost 4 philo
 4. Allow pick chopstik in crit sectn

Algorithm : do {

```
    wait ( chopstick [i] );
    wait ( " [ (i+1)%.5 ] "
           || eat ||
           signal ( ) )
    Signal ( ) -|| think ||

```

} while (true)

• Monitors

- A high lvl sync tool
- Allows only 1 proc to execute at a time
- provides convinient way to manage shared vars

eg monitor abc {
procedure P1 () {
 Pn ()
 some code ()
}
}

• Condition vars

- $x.wait()$ & $x.signal()$
- When a process signals a condition var, it can either :-
 1. signal & wail
 2. signal & continue

- Monitor sol" to dining philo

```
monitor phil {
```

```
    enum { think, hung, eat } state [5];  
    condition self [5];
```

```
void pickup (int i) {
```

```
    state [i] = hung;  
    test (i);
```

```
    if (state [i] != eat) self [i]. wait;
```

```
void putdown (int i) {
```

```
    state [i] = think;
```

```
    test ((i+4)%5);
```

```
    test ((i+1)%5);
```

* dk if imp

```
void test (int i) {
```

```
    if (state [(i+4)%5] != eat && state [i] == hung && state [(i+1)%5]  
        != eat)
```

```
{
```

```
    state [i] = eat;
```

```
    self [i]. signal();
```

```
}
```

```
}
```

```
void initial () {
```

```
    for (int i=0; i<5; i++) {
```

```
        state [i] = think;
```

```
}
```

```
}
```