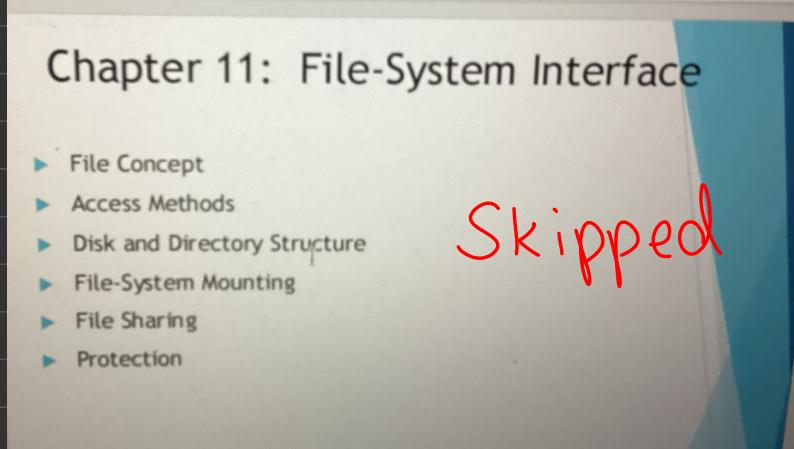


- Topics Covered

1. Swapping
2. Contiguous allocation / Fragmentation / Segmentation
3. Paging
4. Demand Paging (Page fault)
5. Page Replacement
6. Thrashing
7. Disk Structure / Attach / Schedule / Manage
8. RAID



## • Swapping

A process is moved from main memory to temp storage (Backing store) & brought again to come running is swapping

This method allows the system to handle more processes than actually it can handle

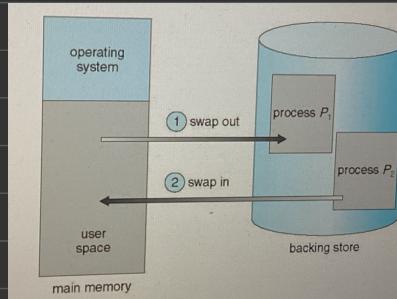
Backing Store : A storage fast enough to hold memory copies of all processes

Roll out, roll in : An algo where low priority processes kicked out to replace high ones

Transfer time : Time Taken to transfer data, depends on size of memory

Pending I/O : If a process has pending I/O it complicates swapping

Need not necessary process comes to same add after swap



## • Context Switching with swapping

When the process in next memory isn't in memory than the OS swaps out one process & swaps in another

Process with pending I/O can't be swapped

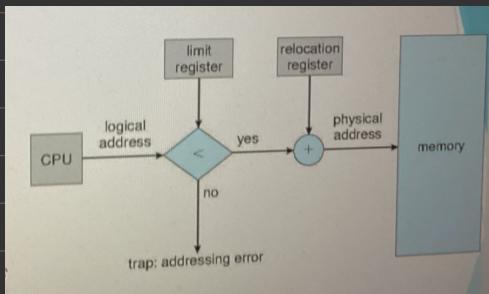
Double buffering : To avoid data issue, it transfers thru kernel space

e.g. : A 1000 mb prcs swapped at 50 mb/s takes 2000 ms to swap out & 2000 ms to swap another

- Contiguous Allocation / Fragmentation / Segmentation

- Contiguous Allocation

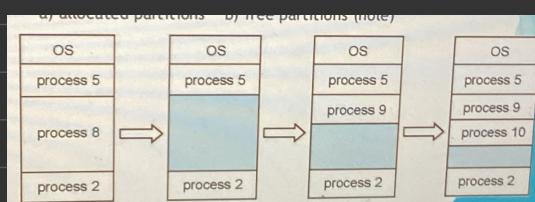
- A basic way of allocating memory by placing each process in a single contiguous block
- Two - memory sections : One for OS (low memory)  
One for user processor (high memory)
- Reallocation reg : A process with reallocated reg is used to interfere of procs
- Base & limit reg : define allowed memory range for each process



- Multiple partition allocation

Here, memory divided into multiple partitions to support multiple process

Memory holes are unused spaces , filled by new processes



In order to allocate memory efficient , there are 3 methods

1. First-fit : First available hole
2. Best-fit : Smallest hole that fits
3. Worst-fit : largest hole

## • Fragmentation

External - Free mmry exists but isn't contiguous causing wasting of space

Internal - Allocated mmry is slightly more than needed, creating unused space

To remove ext. frag we move main mmry & create one large free block

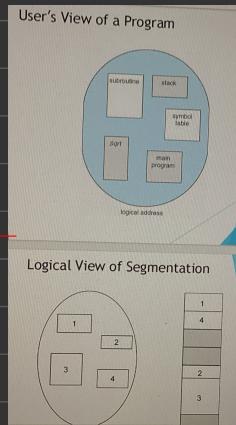
Compaction : Shuffle mmry elems to create free parts, but this can't be done dynamically

## • Segmentation

Organises mmry in segments based on logical divisions

Allows user to see mmry as separate logical units such as :-

- main prog
- func
- method
- obj
- stack
- arrays



### Architecture

- logical add : Contains segment no & offset
- segment table : maps logical add to physical add

Base - contains the start add of pcess  
limits - specifies length

Segment table Base reg (STBR) : points to segment table's loc

" " length reg (STLR) : indicates no. of segments

Protection bits are associated with segments

each entry gives a bit

- Paging

**Non-contiguous allocation** : A process's memory can be spread across diff areas in physical mmry so it doesn't have to be continuous

**Benefits** : Prevent gaps  
Solve issues with diff sizes

**Concepts** : Mem divided into equal part called frames  
Each page fits into a frame

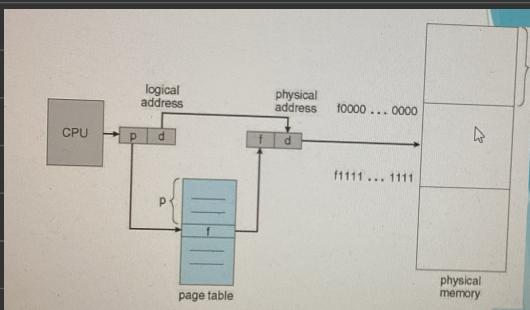
**limit** : Can still lead to wastage of space (int. frag)

- Add translation

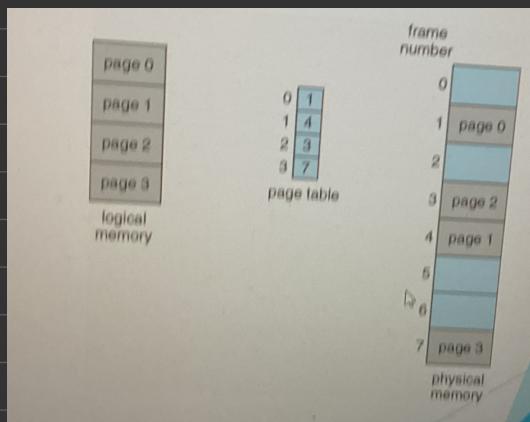
Page number ( $p$ ) : Identifies which page to refer

Page offset ( $d$ ) : Specifies exact location

Paging b/w :



Paging model :



## Implementation of page table

- Stored in main memory

PTBR (page table base reg) : points to table  
 PTLR (length) : indicate size

- Each access req, 2 lookups, one for page & one for data

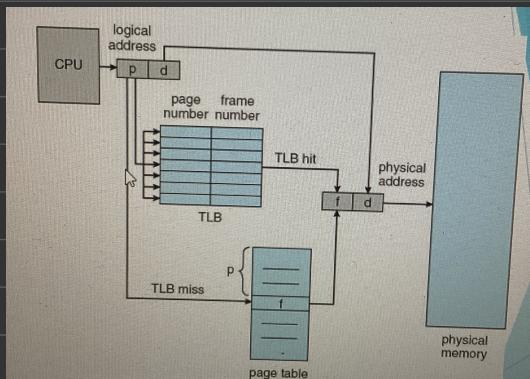
- Solution : TLB (Translation look-aside buffer)

Gives a fast-lookup

### Features

- Some TLB store identifier to manage multiple process
- Usually small
- On a miss, entry used in future
- Replace old entries

## Paging hit with TLB :



## Effective Access time

TLB takes minimal time to search but may not always find page

## Formula on TLB (hit /miss)

$$EAT = (1 + \varepsilon) \alpha + (2 + \varepsilon) (1 - \alpha)$$

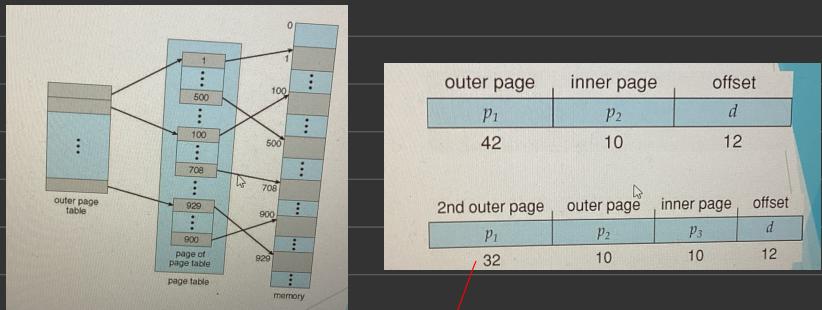
|
hit
miss

## • Memory protection

- Each frame has a bit indicating the process to read / write
- Bit may give read / write only permissions

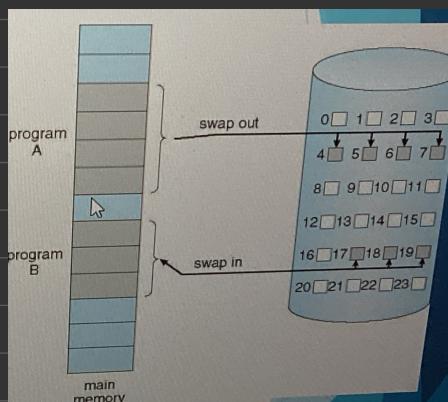
## • More types of Paging

1. Hierarchical Paging
2. Two-lvl paging
3. Hashed paging
4. Three-lvl paging
5. Inverted paging



## • Demand Paging (Page fault)

- Loads only necessary pages of a prog as they're needed , rather than loading the entire prog
- This approach saves mmry & I/O
- Also improves response time & supports users
- When a page isn't in mmry , the OS loads it ( pager )
- Pager predicts which pages needed & only brings those
- If page isn't in mmry , it's load from storage

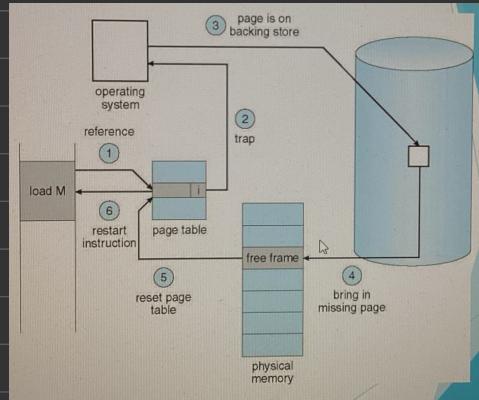


## • Page fault

↳ occurs when page needed by a prog isn't in mmry

1. OS checks if it's valid
2. find a free frame
3. loads page in frame
4. Updates tables
5. Restart Instruc

- Steps :
1. reference
  2. trap
  3. page in back store
  4. missing page
  5. reset page table
  6. Restart instruc



## • Aspects of demand paging

- Pure DP mean no page initially in mmry
- Each instruc fetch can lead to multiple page fits
- HW support req for DP
- Performance of DP
  - Involves 3 main activities : handle interrupt , read page & restart process
  - The page fault rate ( $\rho$ ) how often it occurs (0.1)
  - EAT combines Time & Page fault overhead

## • Page Replacement

**Purpose :** Prevents excessive memory allocation by adding a page-replacement mechanism in the system

**Modify (Dirty) Bit :** Helps reduce the need for data transfers

**Benefit :** Supports large virtual memory by enabling more efficient use

without PR, system will run out of memory quickly

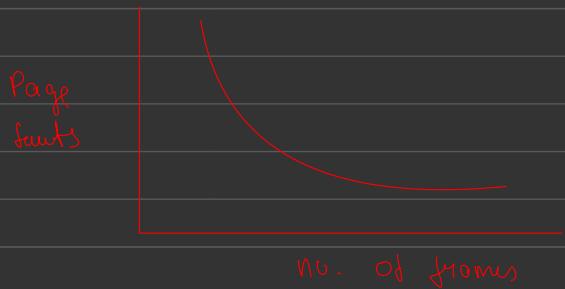
## • Page & Frame replacement algorithms

**Frame-allocation algo :** determines no. of frames assigned to each process & which " " to replace

**Page - Replacement algo :** Aims for lowest page-fault rate by tracking seq.

The algos are :-

1. FIFO
2. Optimal
3. LRU
4. Enhanced 2nd Chance
5. Counting
6. Page buffering



## • FIFO (First In First Out)

**Mech :** Replaces the oldest loaded page

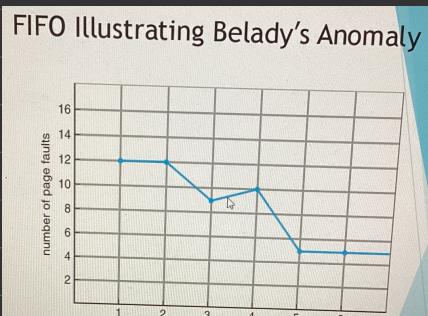
Adding More frames can cause more page faults

It's used to track order of pages

3 frames (3 pages can be in memory at a time per process)		
reference string		
7	0	1
2	0	3
0	4	2
3	0	2
3	2	1
2	1	0
0	3	1
1	2	0
7	0	1
7	7	7
0	0	0
1	1	1

page frames

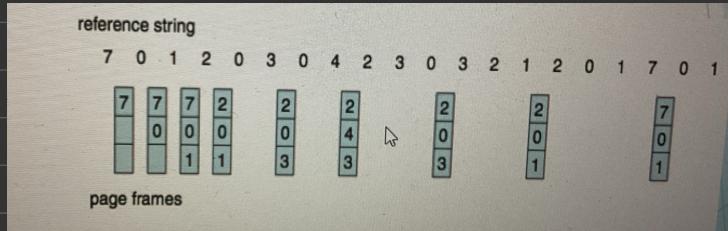
15 page faults



## • Optimal Algorithm

Mech : Replaces the page that won't be used for the longest future period

limit : theoretical since relying on future



## • Last Recently Used (LRU)

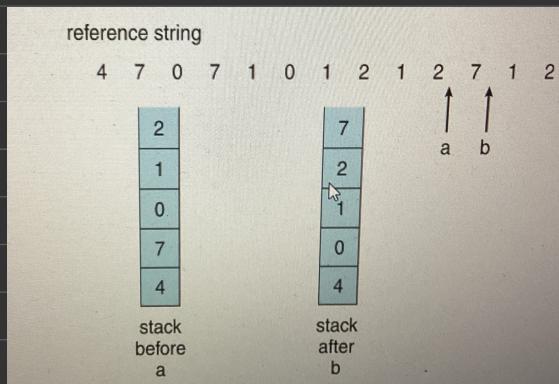
Mech : Replaces the page that has not been used for the longest time

Use : Can use counters / stack to record time

Adv : Generally, effective but complex to implement

Counter Implementation : Each page entry has a counter, increments the clock & record time

Stack " " : Every page is referred to item in stack from Top , Top increments recording time



## • Enhance Second Chance

Uses both the reference & modify bits to classify pages

(0, 0) : neither recently used nor modified

(0, 1) : " " " but "

(1, 0) : recently used not modified

(1, 1) : both

## • Counting Algorithm

Least frequently used (LFU) : Replaces page with fewest responses

Most " " (MFU) : " " " most "

## • Page - Buffering

Mech : Maintains a pool of free frames to speed up page replacement

Pages might remain in free frame until they're needed again

## • Applications

1. Guessing future pg access
2. Better knowledge
3. Save response time
4. Bypass buffering
5. Prevent locking

Memory Full ⓘ

### Allocation of Frames

- Each process requires a minimum number of frames.
- Fixed Allocation: Frames are equally divided among processes.
- Proportional Allocation: Frames are distributed based on process size, which may change dynamically.

### Global vs. Local Allocation

- Global Replacement: Frames are shared among processes, which increases overall system efficiency but causes variation in individual process performance.
- Local Replacement: Limits each process to its allocated frames, providing consistent performance but risking underutilized memory.

### Non-Uniform Memory Access (NUMA)

- NUMA Systems: Access speed varies based on memory location.
- Optimization: Memory and CPU allocations are optimized to improve performance by allocating resources close to the CPU running the process.

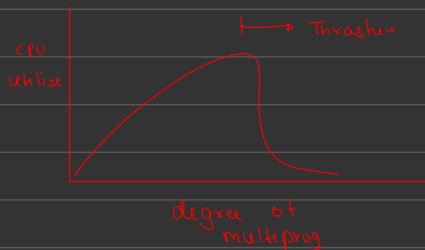
- Thrashing

Thrashing happens when a process doesn't have enough pages causing a high page-fault rate

It leads to constant swapping of pages in & out of memory because pages that were just removed again are quickly needed again

This Causes :-

- Low CPU utilization
- OS mistakes making worse problems
- More processes being added, inc thrashing further



- Demand Paging & thrashing

- DP works based on locality model
- Thrashing occurs when the size of these locality models exceed
- To reduce thrashing, we use local / priority based page replacement strategies

- Working Set Model

- The " " " ( $\Delta$ ) is a recent set of page references  
eg last 10000 instruc
- Working Set (WSS) of a process = Total pages accessed

if  $\Delta$  too small : may not capture full set of pages

$\Delta$  too large : may include pages not currently in use

$\Delta = \infty$  : It includes entire proc

$D = \Sigma (wss)$  : represents total demand

\* Thrashing occurs if  $D$  exceeds the available memory frames

#### Slide 5: Keeping Track of the Working Set

- Track working sets using an interval timer and a reference bit.
  - Example:  $\Delta = 10,000$ ; timer interrupts every 5,000 time units.
  - Two bits per page are used to check if the page has been referenced recently.
- This approach approximates working sets but may not be completely accurate.
- Improvement: Use more bits (e.g., 10 bits) and shorter intervals for better tracking.

#### Slide 6: Page-Fault Frequency (PFF)

- PFF is a straightforward approach to manage page faults by setting an acceptable page-fault frequency rate.
- Uses a local replacement policy:
  - If a process's page-fault rate is too low, it loses a frame (indicating it doesn't need more memory).
  - If the rate is too high, it gains a frame (indicating it needs more memory).

## • Disk Structure / Attachment / Scheduling / Management

### • Disk Structure

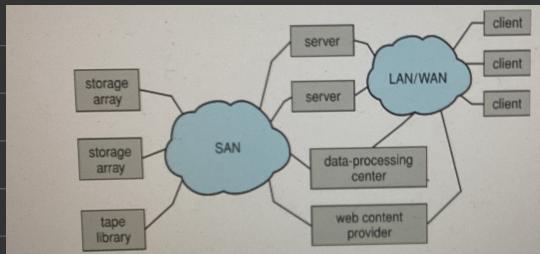
- Disk drives are organised as large 1D arrays of small storage unit logical block
- Formatting sets up these logical blocks
- Disk Structure : starts from outermost cylinder & moves inward
- Mapping logical add to phy add is easy but issue can arise with bad sectors
- Disks may have diff no. of sectors per track

## • Disk Attachment

- Disks connect to comps via I/O ports & buses
- **SCSI** : A type of bus that can connect upto 16 devices at a time
- SCSI devices have initiators (start tasks) & targets (perform tasks)
- **Fibre Channel** : A high speed connection used in SAN (Storage area network) to link multiple Storage devices

SAN :

- Its one /more storage arrays
- Hosts also attach to switch
- Multiple hosts " to multiple array



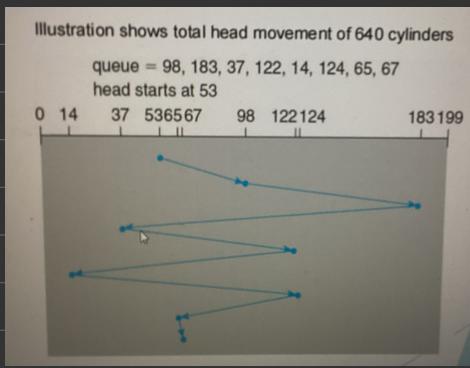
## • Disk Scheduling

- OS manages disk to ensure quick access & efficient use
- Minimize seek time
- The OS schedule disk task to minimize delays
- Disk I/O req from :
  - OS
  - System prcs
  - User prcs
- Types of Disk scheduling Algo
  - 1. FCFS
  - 2. SSTF
  - 3. SCAN
  - 4. C-SCAN
  - 5. C-LOOK
- Diff algo are used to decide best order

## 1. FCFS

(First come, first Serve)

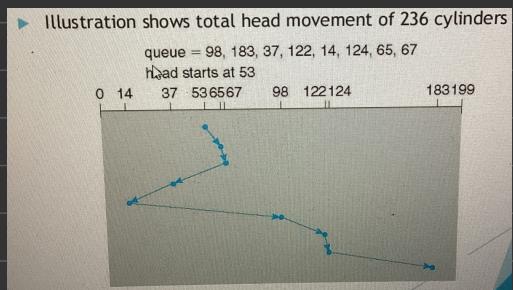
Processes req. in order, but res in unnecessary head movement



## 2. SSTF

(Shortest seek time first)

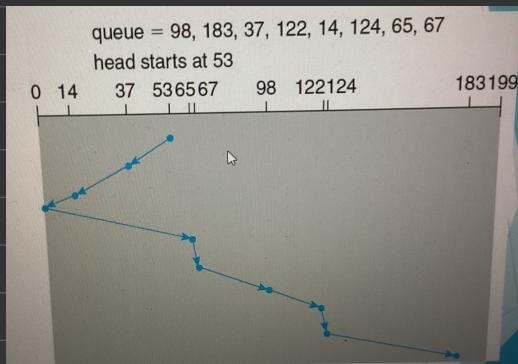
- Chooses the closest req. to minimize head movement
- Some req. may wait longer
- Commonly used n simple



## 3. SCAN

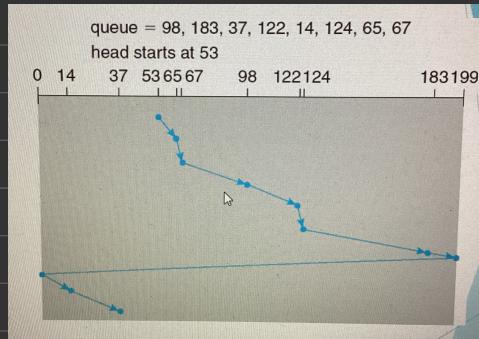
(Elevator Algo)

- The head moves across the disk
- Serving req. in one direction then reverse
- Used in high-demand disk



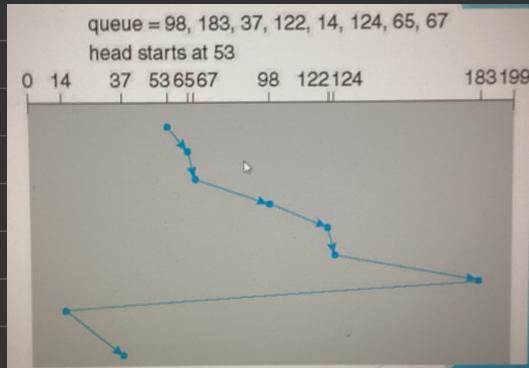
#### 4. C-SCAN

- Moves in one direction
- Once reached the end, returns w/o servicing req
- Provides fair WT
- Used in high-demand disk



#### 5. C-LOOK

- Like LOOK a version of SCAN, C-LOOK version of C-SCAN
- Go to the last req. in each direction



#### Disk Management

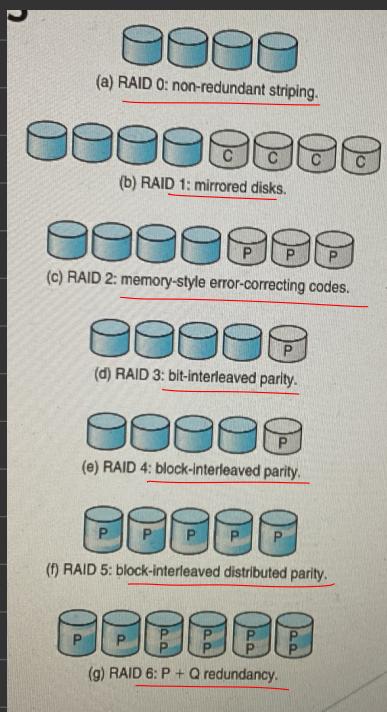
- low-level formatting prepares the disk by dividing it into **sectors**, store units
- **Partitioning** organises the disk into sections
- **logical formatting** creates a file system structure
- **RAW disk access** allows apps like db to manage their own disk

- Raid Structure

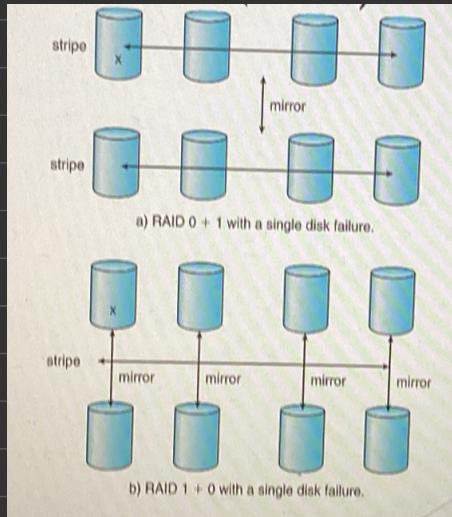
- RAID (Redundant Array of inexpensive disk) is a system that uses multiple disks together
- To improve reliability & data redundancy
- Using multiple disks means longer time before a failure (mean time to failure)
- Mean time to repair (MTTR): It's a time taken to fix a failed disk
  - During this time, another failure could risk data
  - Mean time to data loss: It depends on foll factors
    - If disk independant
    - long mean time to failure
- disk striping: Split data across multiple disks, treating them as a single unit
- Raid lvls: Raid has six lvls with different ways of organizing data for speed
- Redundant data storage: This setup improves performance & safety
  1. Mirror /Shadowing (RAID 1): Keep exact copy of each disk
  2. Stripped Mirror (RAID 1+0): Combine speed & reliability
  3. Block interleaved (RAID 4,5,6): Uses fewer backups
- RAID can still fail if storage array fails, so data replication is common

- RAID lvls

RAID  
1  
2  
3  
4  
5  
6



- Raid (0+1) & (1+0)



- Other features

1. Snapshot : Capture the state of file at a specific time , helpful for backup

2. Replication : Automatically duplicates data to other locations for backup

3. Hot Spare Disk : An used disk that automatically takes over if another fails